

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>




deeplearning.ai



Activations (Basic Properties)

Outline

- What are activations
- Reasoning behind non-linear differential activations



Activations



$$z_i^{[l]} = \sum_{i=0} W_i^{[l]} a_i^{[l-1]}$$

$$a_i^{[l]} = \boxed{g^{[l]}}(z_i^{[l]})$$

Differentiable
non-linear
function


Activations

$$a_i^{[l]} = \boxed{g^{[l]}}(z_i^{[l]})$$

Differentiable
non-linear
function

1. Differentiable for backpropagation

2. Non-linear to compute complex features, **if not**:



\equiv

$$WX + b$$

Linear
regression

Summary

- Activation functions are non-linear and differentiable
- Differentiable for backpropagation
- Non-linear to approximate complex functions





deeplearning.ai

Common Activation Functions

Outline


- Common activations and their structure
 - ReLU
 - Leaky ReLU
 - Sigmoid
 - Tanh



Activations: ReLU


ReLU = Rectified Linear
Unit

$$g^{[l]}(z^{[l]}) = \max(\underline{0}, \underline{z^{[l]}})$$



Dying ReLU
problem

Activations: Leaky ReLU




$$g^{[l]}(z^{[l]}) = \max(\underline{az^{[l]}}, \underline{z^{[l]}})$$

Solves the dying
ReLU problem

Activations: Sigmoid

$$g^{[l]}(z^{[l]}) = \frac{1}{1 + e^{-z^{[l]}}}$$




Values between 0
and 1

Vanishing gradient
and saturation
problems

Activations: Tanh

$$g^{[l]}(z^{[l]}) = \tanh(z^{[l]})$$



Values between -1
and 1

Keeps the sign of
the input

Same issues as
Sigmoid

Summary

- ReLU activations suffer from dying ReLU
- Leaky ReLU solve the dying ReLU problem
- Sigmoid and Tanh have vanishing gradient and saturation problems





deeplearning.ai

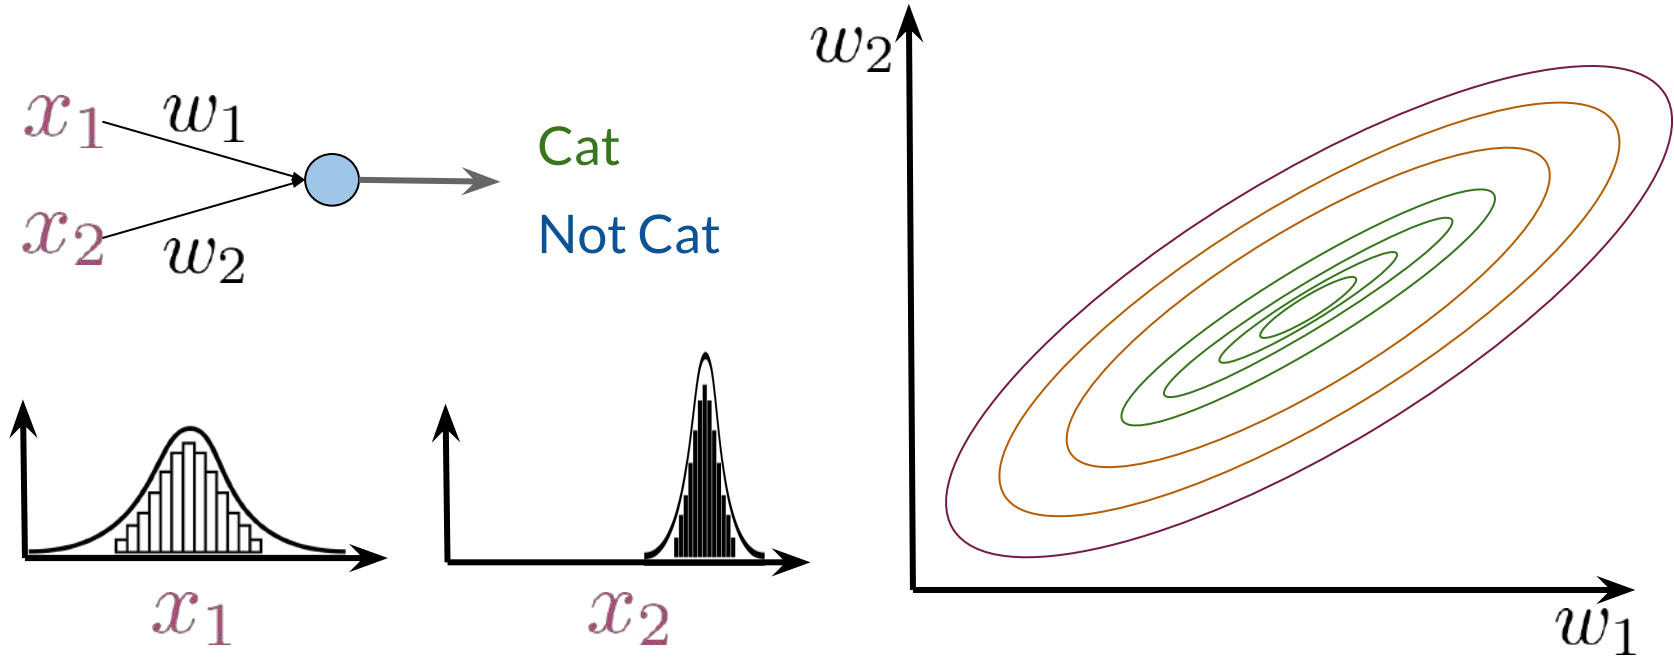
Batch Normalization (Explained)

Outline


- How normalization helps models
- Internal covariate shift
- Batch normalization



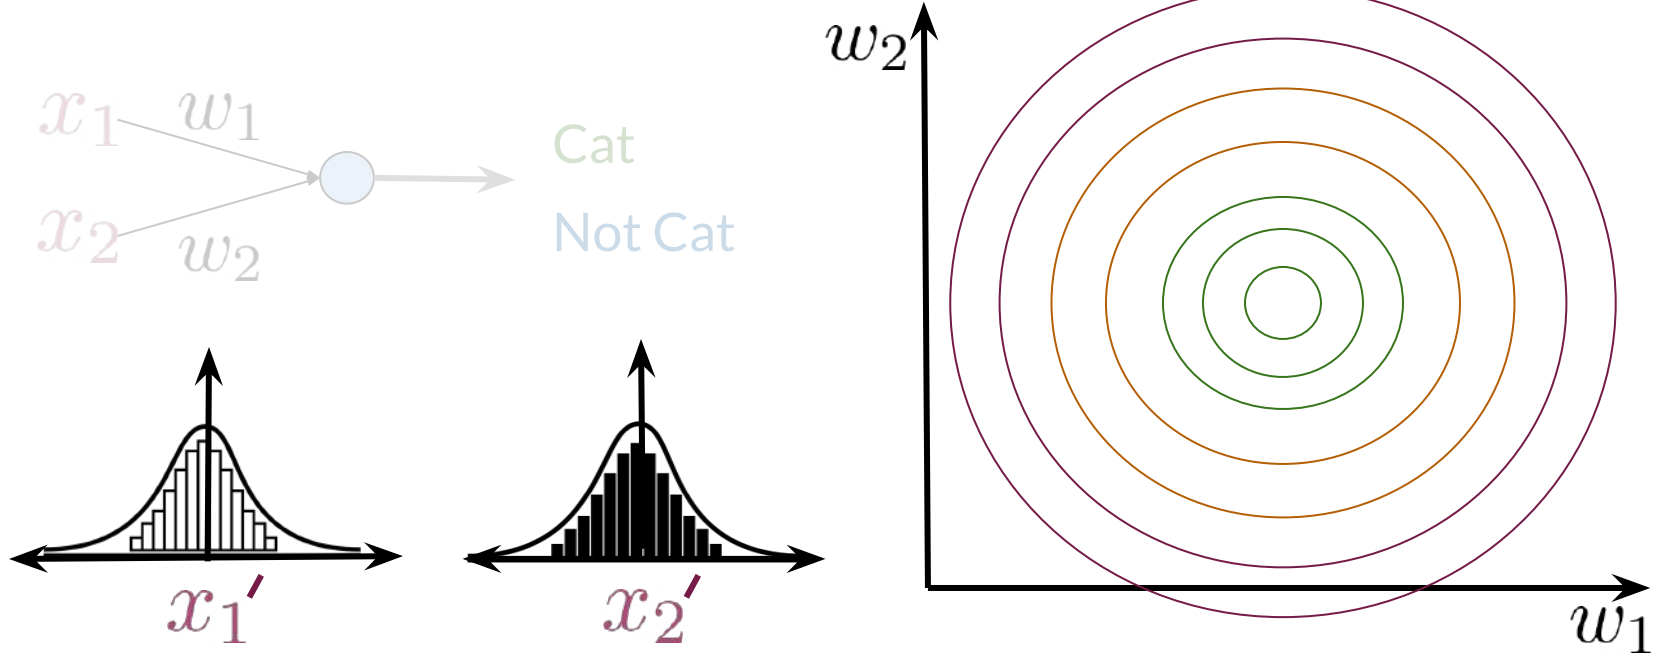
Different Distributions




Covariate Shift



Normalization and Its Effects



Normalization and Its Effects




x_1' x_2'

Around mean at 0
and std. at 1

Training data uses
batch stats

Test data uses
training stats

Normalization and Its Effects




x_1' x_2'

Around mean at 0
and std. at 1

Reduction of
covariate shift

Internal Covariate Shift




Changes in
weights




Changes in
activation
distribution

Batch Normalization




Normalizes the
input for each
neuron



Summary

- Batch normalization smooths the cost function
- Batch normalization reduces the internal covariate shift
- Batch normalization speeds up learning!






deeplearning.ai



Batch Normalization (Procedure)

Outline

- Batch norm for training
- Batch norm for testing



Batch Normalization: Training





$$z_i^{[l]} = \sum_{i=0} W_i^{[l]} a_i^{[l-1]} \rightarrow \text{For every example in the batch}$$

$$\hat{z}_i^{[l]} = \frac{z_i^{[l]} - \mu_{z_i^{[l]}}}{\sqrt{\sigma_{z_i^{[l]}}^2 + \epsilon}}$$

Batch mean of $z_i^{[l]}$
Batch std of $z_i^{[l]}$

Batch Normalization: Training



$$\hat{z}_i^{[l]} = \frac{z_i^{[l]} - \mu_{z_i^{[l]}}}{\sqrt{\sigma_{z_i^{[l]}}^2 + \epsilon}}$$


$$y_i^{[l]} = \boxed{\gamma} \hat{z}_i^{[l]} + \boxed{\beta}$$

Shift factor

Scale Factor

Learnable
parameters to get
the optimal dist.

Batch Normalization: Test



$$\hat{z}_i^{[l]} = \frac{z_i^{[l]} - \mathbf{E}(z_i^{[l]})}{\sqrt{\text{Var}(z_i^{[l]}) + \epsilon}}$$

Running **mean** and running **std** from training

Frameworks like
Tensorflow and Pytorch
keep track of them

Summary

- Batch norm introduces learnable shift and scale factors
- During test, the running statistics from training are used
- Frameworks take care of the whole process






deeplearning.ai


Review of Convolutions

Outline


- What convolutions are
- How they work




What is a convolution?




What is a convolution?



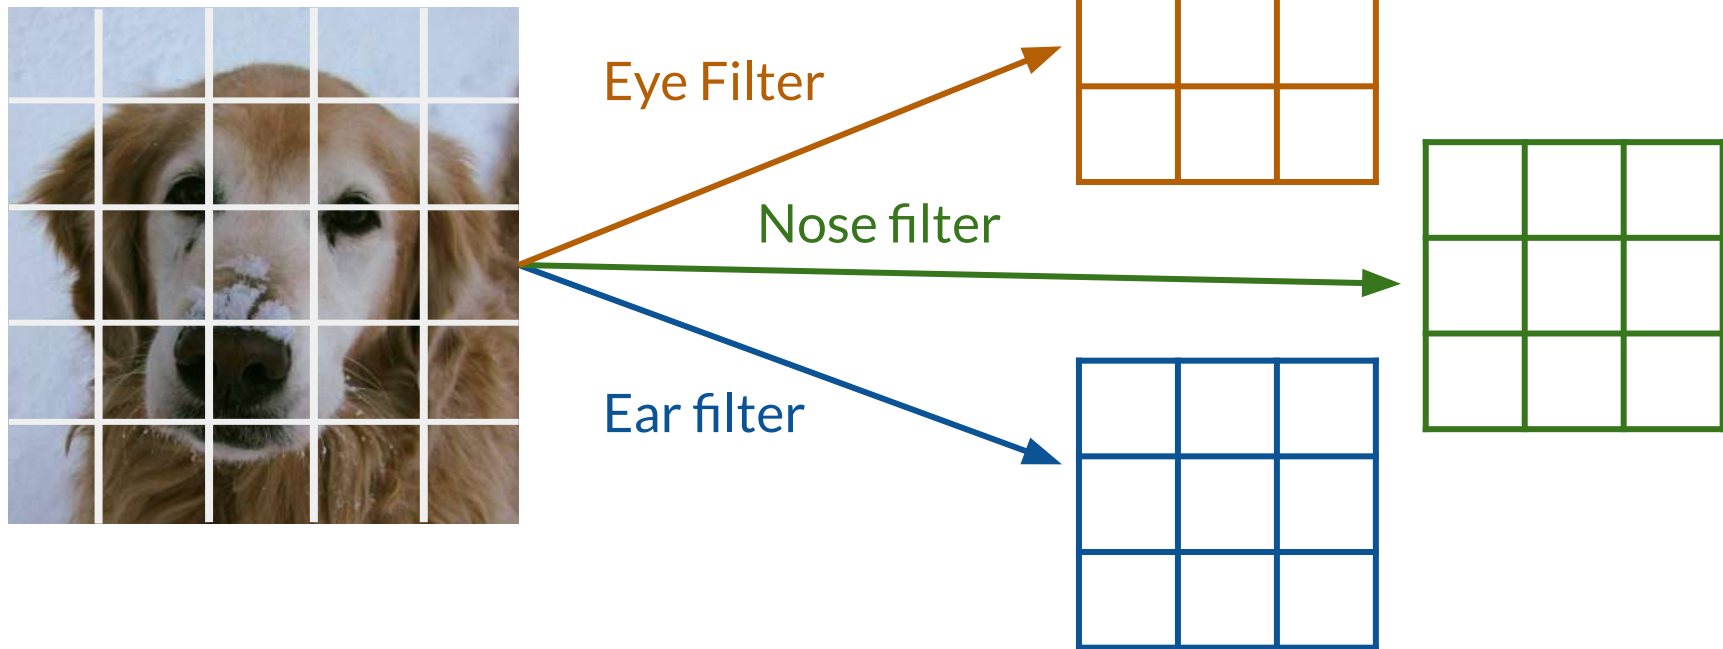
Eye Filter




What is a convolution?



What is a convolution?



What is a convolution?



Eye Filter

5	3	1
10	23	1
1	42	4

Nose filter

2	21	5
23	45	12
3	32	4

Ear filter

32	2	24
25	12	66
3	45	2

What is a convolution?

50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image

0 (black) to 255 (white)

What is a convolution?

50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

0 (black) to 255 (white)

What is a convolution?

50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image



Convolution

Filter

1	0	-1
1	0	-1
1	0	-1

0 (black) to 255 (white)

What is a convolution?

50×1	50×0	0×-1	0	0
50×1	50×0	0×-1	0	0
50×1	50×0	0×-1	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image



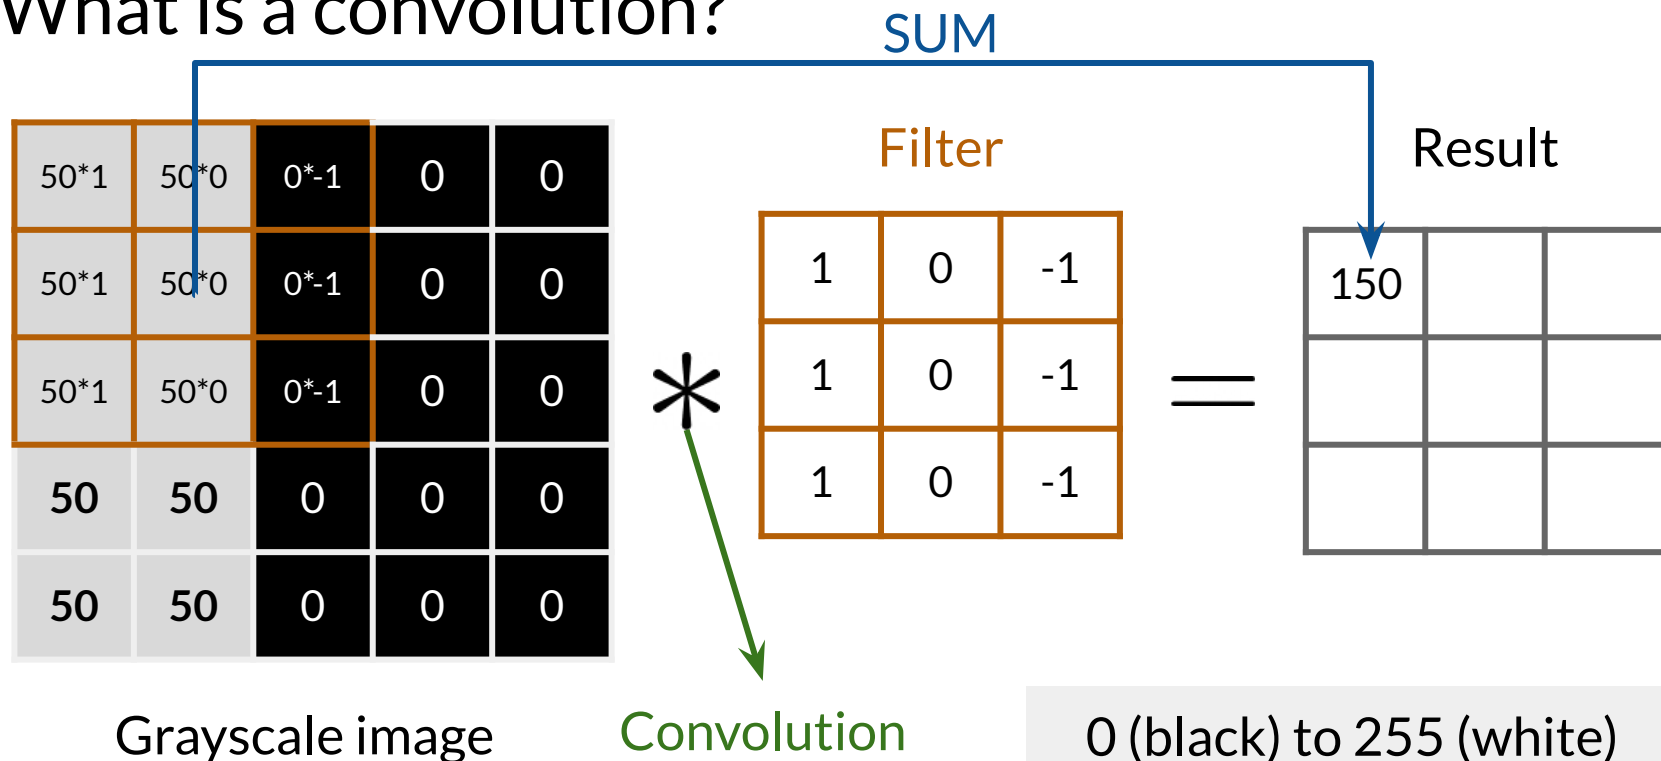
Convolution

Filter


1	0	-1
1	0	-1
1	0	-1

0 (black) to 255 (white)


What is a convolution?




What is a convolution?




What is a convolution?




What is a convolution?



What is a convolution?



What is a convolution?



Summary

- Convolutions are useful layers for processing images
- They scan the image to detect useful features
- Just element-wise products and sums!






deeplearning.ai

Padding and Stride

Outline

- Padding and stride
- The intuition behind padding



Stride

50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

Stride

50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

=

Result

150		

Stride

→ 1 Pixel to the right

50	50*1	0*0	0*-1	0
50	50*1	0*0	0*-1	0
50	50*1	0*0	0*-1	0
50	50	0	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

=

Result

150	150	

Stride

→ 1 Pixel to the right

50	50	0*1	0*0	0*-1
50	50	0*1	0*0	0*-1
50	50	0*1	0*0	0*-1
50	50	0	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

=

Result

150	150	0

Stride

→ 1 Pixel to the right

↓ 1 Pixel down

50	50	0	0	0
50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

=

Result

150	150	0
150		

Stride

→ 1 Pixel to the right



1 Pixel down

50	50	0	0	0
50	50*1	0*0	0*-1	0
50	50*1	0*0	0*-1	0
50	50*1	0*0	0*-1	0
50	50	0	0	0

Grayscale image



Filter

1	0	-1
1	0	-1
1	0	-1



Result

150	150	0
150	150	

Stride

→ 1 Pixel to the right



1 Pixel down

50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image



Filter

1	0	-1
1	0	-1
1	0	-1




Result

150	150	0
150	150	0
150	150	0

Stride = 1

Stride



Stride

→ 2 Pixels to the right

↓ 2 Pixels down

50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50	50	0	0	0
50	50	0	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1

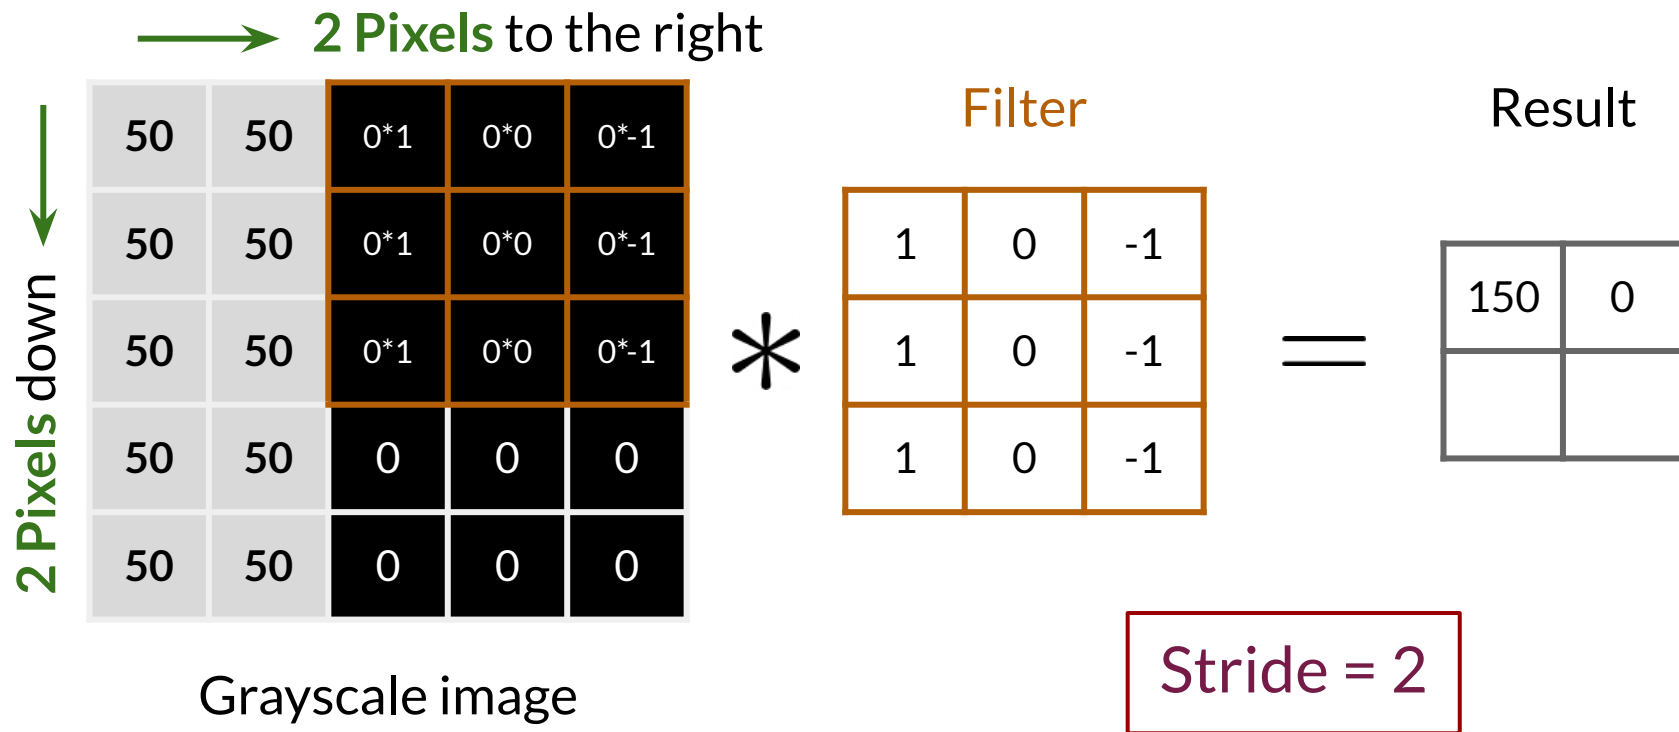
=

Result

150	

Stride = 2

Stride



Stride

→ 2 Pixels to the right

↓ 2 Pixels down

50	50	0	0	0
50	50	0	0	0
50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0
50*1	50*0	0*-1	0	0

Grayscale image

*

Filter

1	0	-1
1	0	-1
1	0	-1


=

Result

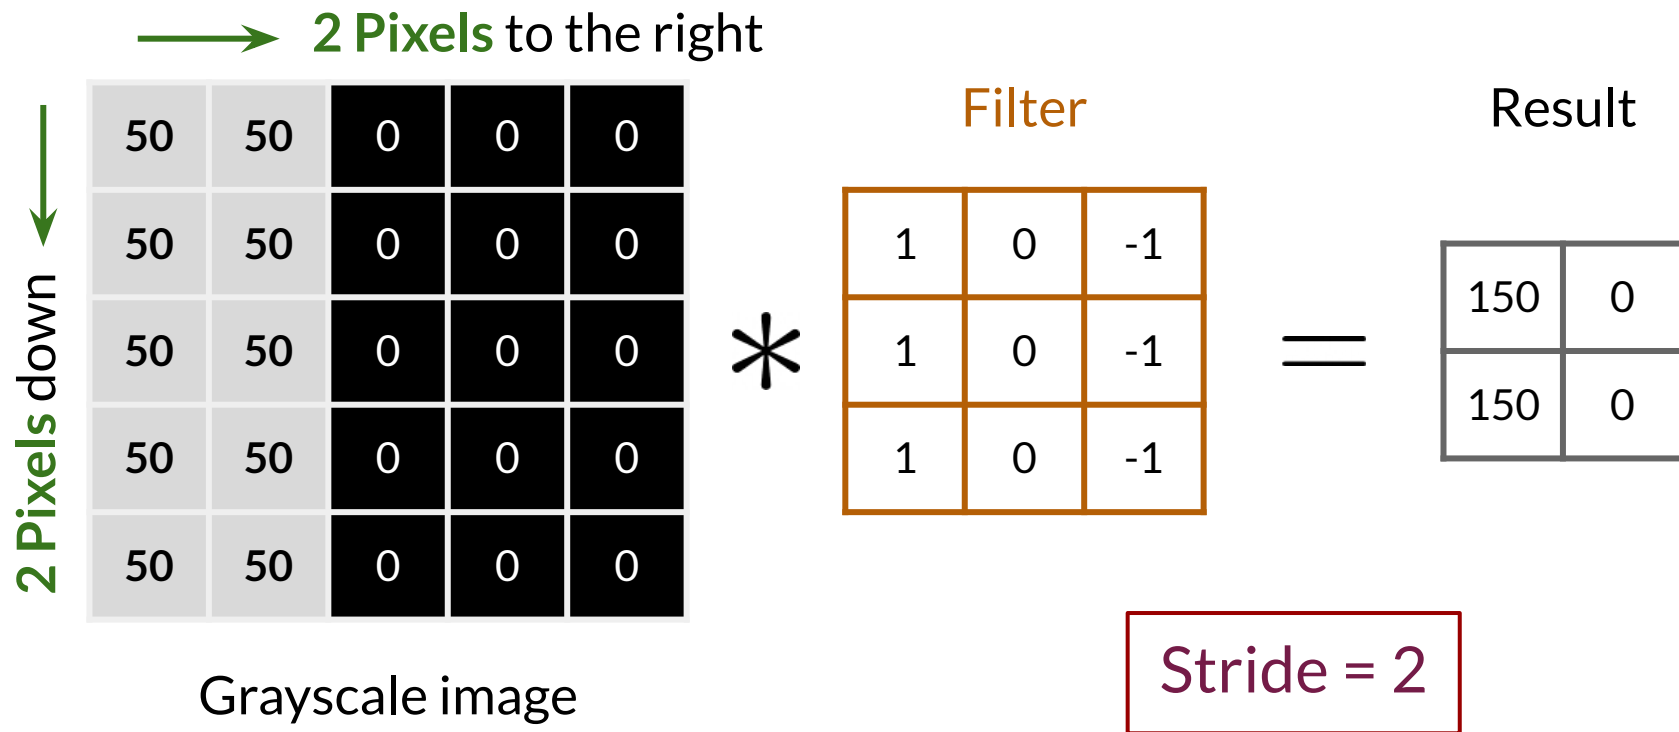
150	0
150	

Stride = 2


Stride




Stride



Padding




Padding



Summary

- Stride determines how the filter scans the image
- Padding is like a frame on the image
- Padding gives similar importance to the edges and the center






deeplearning.ai


Pooling and Upsampling

Outline


- Pooling
- Upsampling and its relation to pooling



Pooling



Max Pooling



4x4 input to 2x2 output

Max Pooling

0	20	30	45
8	1	23	43
0	4	40	20
10	6	43	42

2x2 pooling with stride = 2

Max pooling

4x4 input to 2x2 output

Max Pooling

0	20	30	45
8	1	23	43
0	4	40	20
10	6	43	42


2x2 pooling with stride = 2

Max pooling

20	


4x4 input to 2x2 output

Max Pooling




4x4 input to 2x2 output

Max Pooling



Max Pooling



4x4 input to 2x2 output

Max Pooling

0	20	30	45
8	1	23	43
0	4	40	20
10	6	43	42

2x2 pooling with stride = 2

Max pooling


20	45
10	43

4x4 input to 2x2 output

Other types include:

1. Average pooling
2. Min pooling

Upsampling




Upsampling




Upsampling: Nearest Neighbors

2x2 input to 4x4 output




Upsampling: Nearest Neighbors

2x2 input to 4x4 output



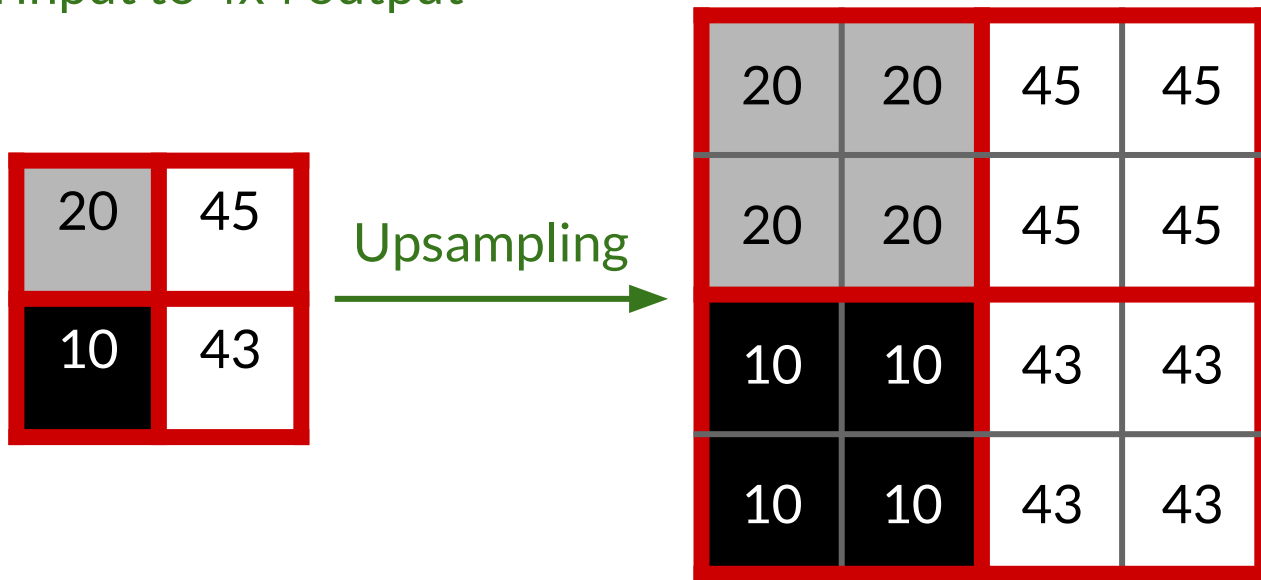
Upsampling: Nearest Neighbors

2x2 input to 4x4 output



Upsampling: Nearest Neighbors

2x2 input to 4x4 output



Upsampling: Nearest Neighbors

2x2 input to 4x4 output

20	45
10	43

Upsampling


20	20	45	45
20	20	45	45
10	10	43	43
10	10	43	43

Other types include:

1. Linear interpolation
2. Bi-linear interpolation

Summary

- Pooling reduces the size of the input
- Upsampling increases the size of the input
- No learnable parameters!






deeplearning.ai

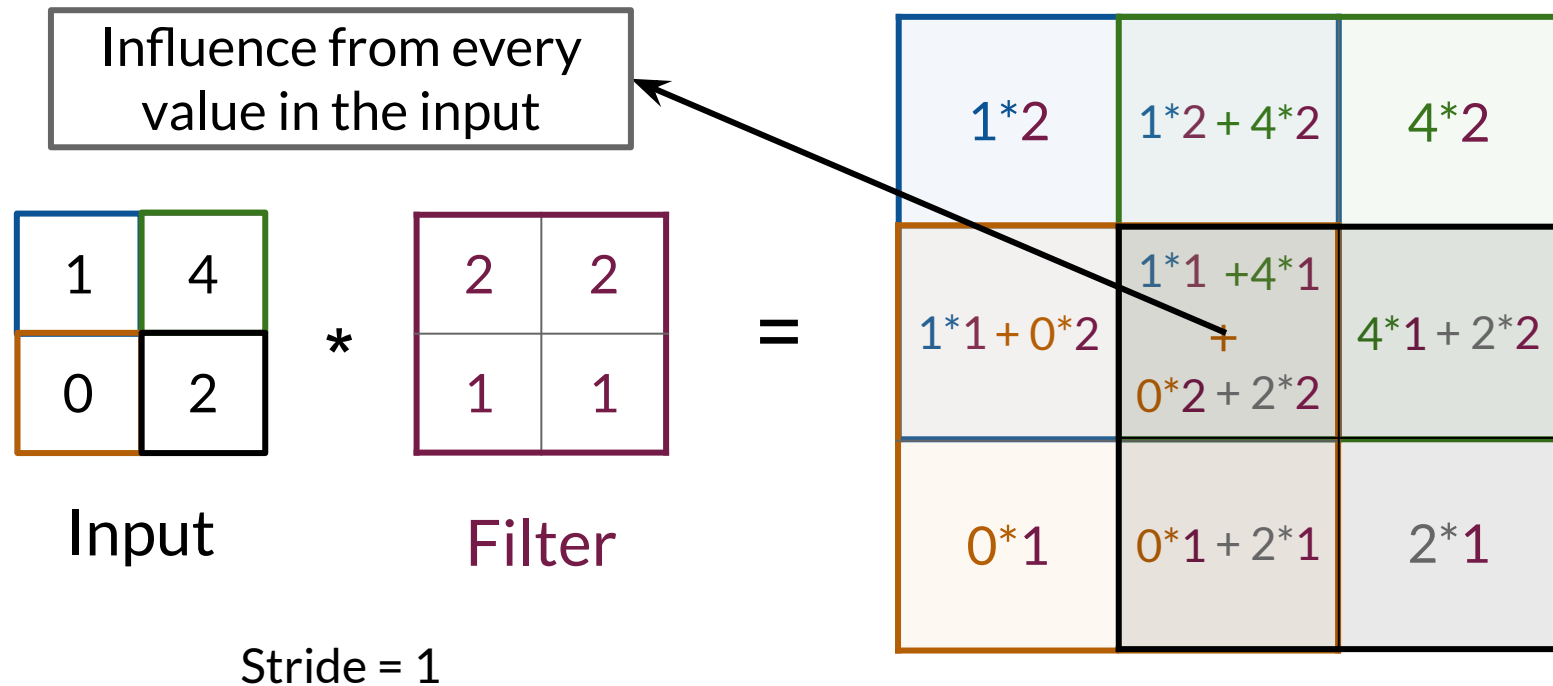
Transposed Convolutions

Outline


- Transposed convolutions as an upsampling technique
- Issues with transposed convolutions





Transposed Convolution



Transposed Convolution



The Problems with Transposed Convolution



Checkerboard
Pattern

Available from: <http://doi.org/10.23915/distill.00003>

Summary

- Transposed convolutions upsample
- They have learnable parameters
- Problem: results have a checkerboard pattern

