

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>



deeplearning.ai

Image-to-Image Translation

Outline

- Image-to-image translation
- Other types of translation






Image-to-Image Translation



Transformation



Available from: <https://twitter.com/citnajt/status/1124904251128406016>

Image-to-Image Translation



Available from: <https://github.com/NVIDIA/pix2pixHD>


Image-to-Image Translation



Available from: <https://youtu.be/3RYNThid23g>

Paired Image-to-Image Translation

Labels to facade



Black-and-white to color



Available from: <https://arxiv.org/abs/1611.07004>

Paired Image-to-Image Translation

Day to night



Edges to photo




Available from: <https://arxiv.org/abs/1611.07004>

Paired Image-to-Image Translation



Clothes and pose to
pose with clothes



Available from: <https://arxiv.org/abs/1705.09368>


Other Translations

“This bird is red with white and has a very short beak”



Available from: <https://arxiv.org/abs/1711.10485>


Other Translations



Available from: <https://arxiv.org/abs/1905.08233>

Summary

- Image-to-image translation transforms images into different styles
- GANs' realistic generation abilities are well-suited to image-to-image translation tasks
- Other types of translation include text-to-image or image-to-video






deeplearning.ai

Pix2Pix Overview

Outline


- Overview of Pix2Pix
- Comparison with conditional GAN
- Upgraded generator and discriminator architectures




Pix2Pix for Paired Image-to-Image Translation

Image-to-Image \longrightarrow *Pix-to-Pix* \longrightarrow *Pix2Pix*

Pix2Pix Generator




Pix2Pix Generator



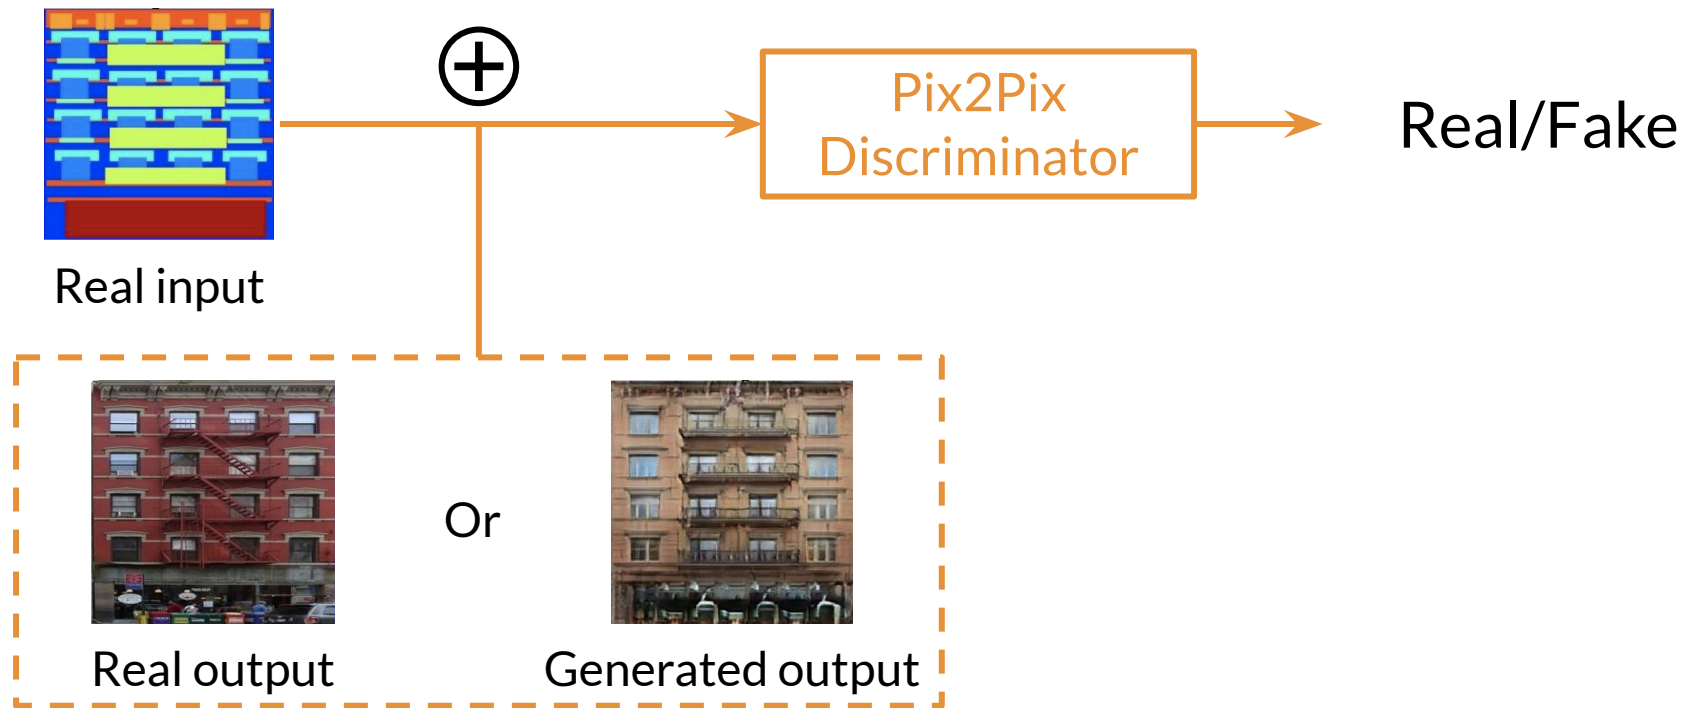
Available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix Generator




Available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix Discriminator





Available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix Upgrades




Based on: <https://arxiv.org/abs/1611.07004>

Pix2Pix Upgrades




(Left) Based on: <https://arxiv.org/abs/1611.07004>
(Right) Based on: <https://arxiv.org/abs/1803.07422>

Pix2Pix Upgrades




Goal is still to produce realistic outputs!



(Left) Based on: <https://arxiv.org/abs/1611.07004>
(Right) Based on: <https://arxiv.org/abs/1803.07422>

Summary

- Inputs and outputs of Pix2Pix are similar to a conditional GAN
 - Take in the original image, instead of the class vector
 - No explicit noise as input
- Generator and discriminator models are upgraded






deeplearning.ai

Pix2Pix: PatchGAN

Outline

- PatchGAN discriminator architecture
- Matrix output vs. single output



Pix2Pix Discriminator: PatchGAN




Image available from: <https://arxiv.org/abs/1611.07004>
Based on: <https://arxiv.org/abs/1803.07422>

Pix2Pix Discriminator: PatchGAN




Image available from: <https://arxiv.org/abs/1611.07004>
Based on: <https://arxiv.org/abs/1803.07422>

Pix2Pix Discriminator: PatchGAN




Image available from: <https://arxiv.org/abs/1611.07004>

Based on: <https://arxiv.org/abs/1803.07422>

Pix2Pix Discriminator: PatchGAN





Image available from: <https://arxiv.org/abs/1611.07004>
Based on: <https://arxiv.org/abs/1803.07422>

Pix2Pix Discriminator: PatchGAN




Still uses BCE Loss

Image available from: <https://arxiv.org/abs/1611.07004>
Based on: <https://arxiv.org/abs/1803.07422>

Pix2Pix Discriminator: PatchGAN

Values closer to *0* = *fake* label




Still uses BCE Loss

Image available from: <https://arxiv.org/abs/1611.07004>
Based on: <https://arxiv.org/abs/1803.07422>

Pix2Pix Discriminator: PatchGAN

Values closer to **1** = **real** label




Still uses BCE Loss

Image available from: <https://arxiv.org/abs/1611.07004>
Based on: <https://arxiv.org/abs/1803.07422>

Summary

- PatchGAN discriminator outputs a matrix of values, each between 0 and 1
- Label matrices:
 - 0's = fake
 - 1's = real





deeplearning.ai

Pix2Pix: U-Net

Outline

- Net framework
 - Encoder-Decoder
- U-Skip connections
- Pix2Pix generator




Image Segmentation




Image Segmentation

Available from: <https://developer.nvidia.com/blog/image-segmentation-using-digits-5/>

Image Segmentation






Image Segmentation

Image-to-Image
Translation




Available from: <https://developer.nvidia.com/blog/image-segmentation-using-digits-5/>

U-Net Framework: Encoder-Decoder




Based on: <https://arxiv.org/abs/1611.07004>

U-Net Framework: Skip Connections



Based on: <https://arxiv.org/abs/1611.07004>

U-Net Framework: Skip Connections



Forward pass


Skip connections
allow information
flow to the decoder

Based on: <https://arxiv.org/abs/1611.07004>

U-Net Framework: Skip Connections


Backward pass

Skip connections
improve gradient
flow to encoder



Based on: <https://arxiv.org/abs/1611.07004>

Pix2Pix Encoder



Pix2Pix Encoder

χ Input size: 256 x 256 x 3




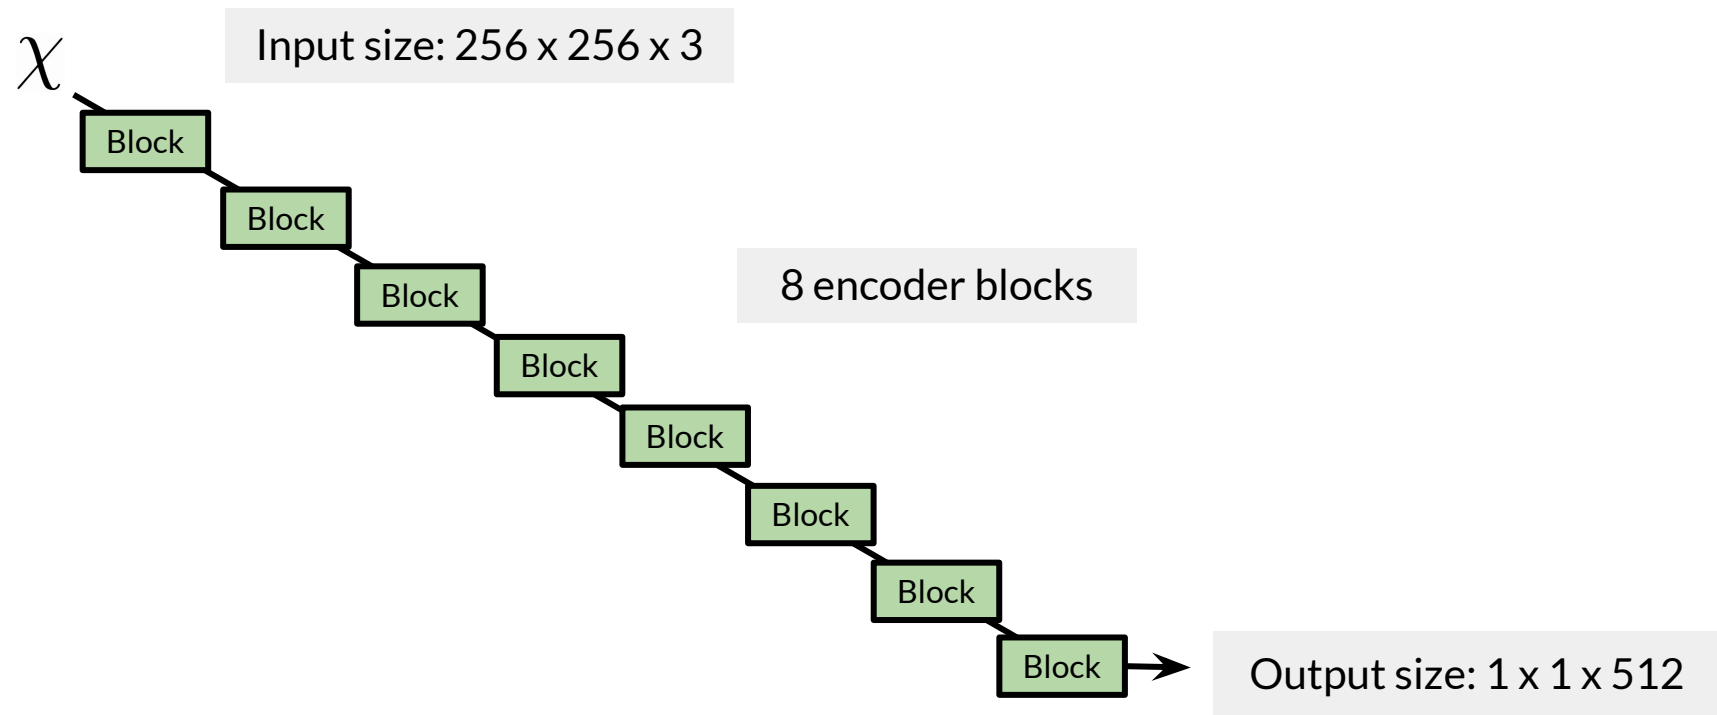


Image available from: <https://arxiv.org/abs/1611.07004>

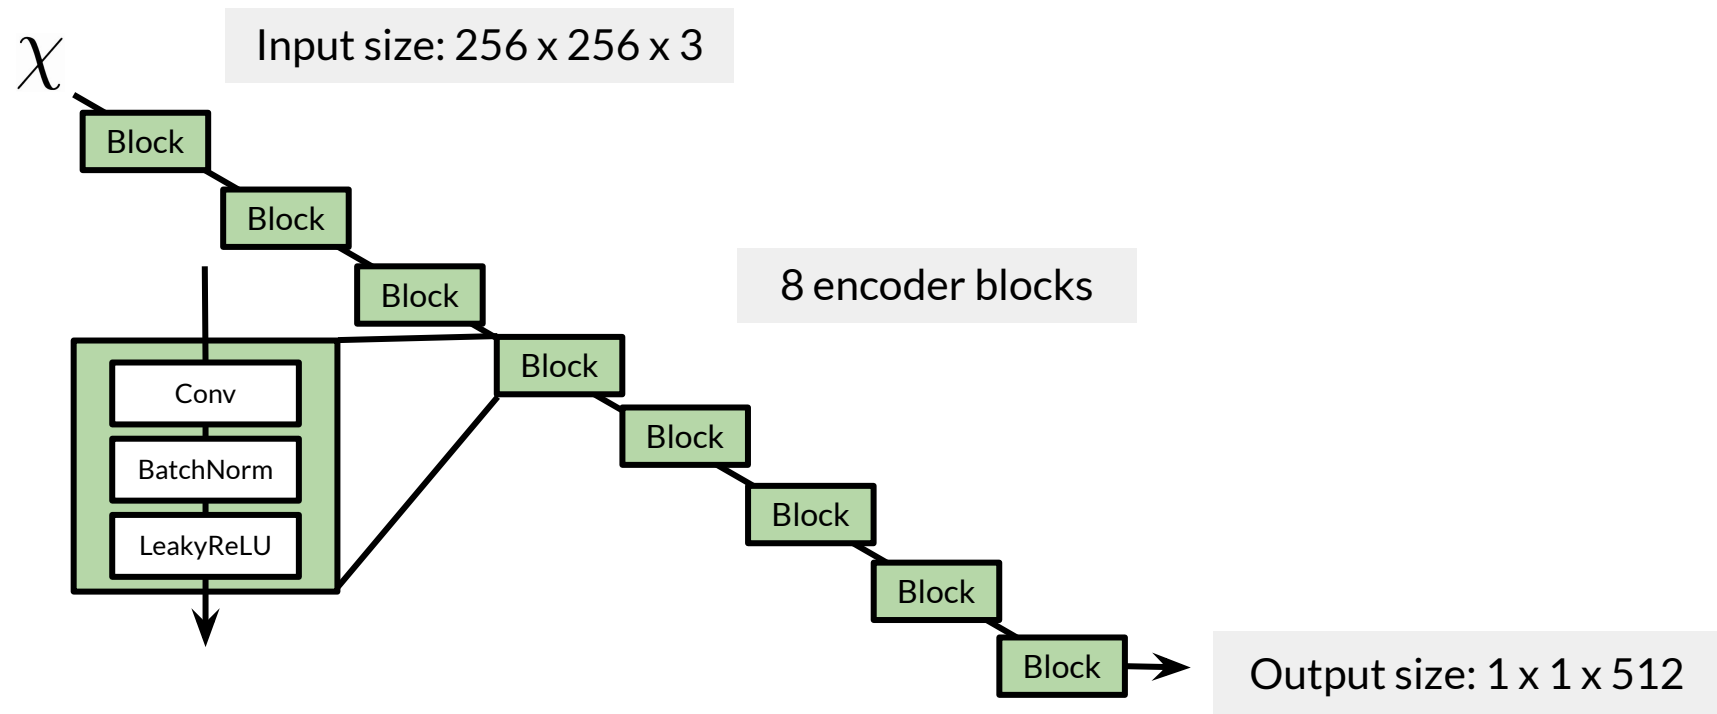
Pix2Pix Encoder



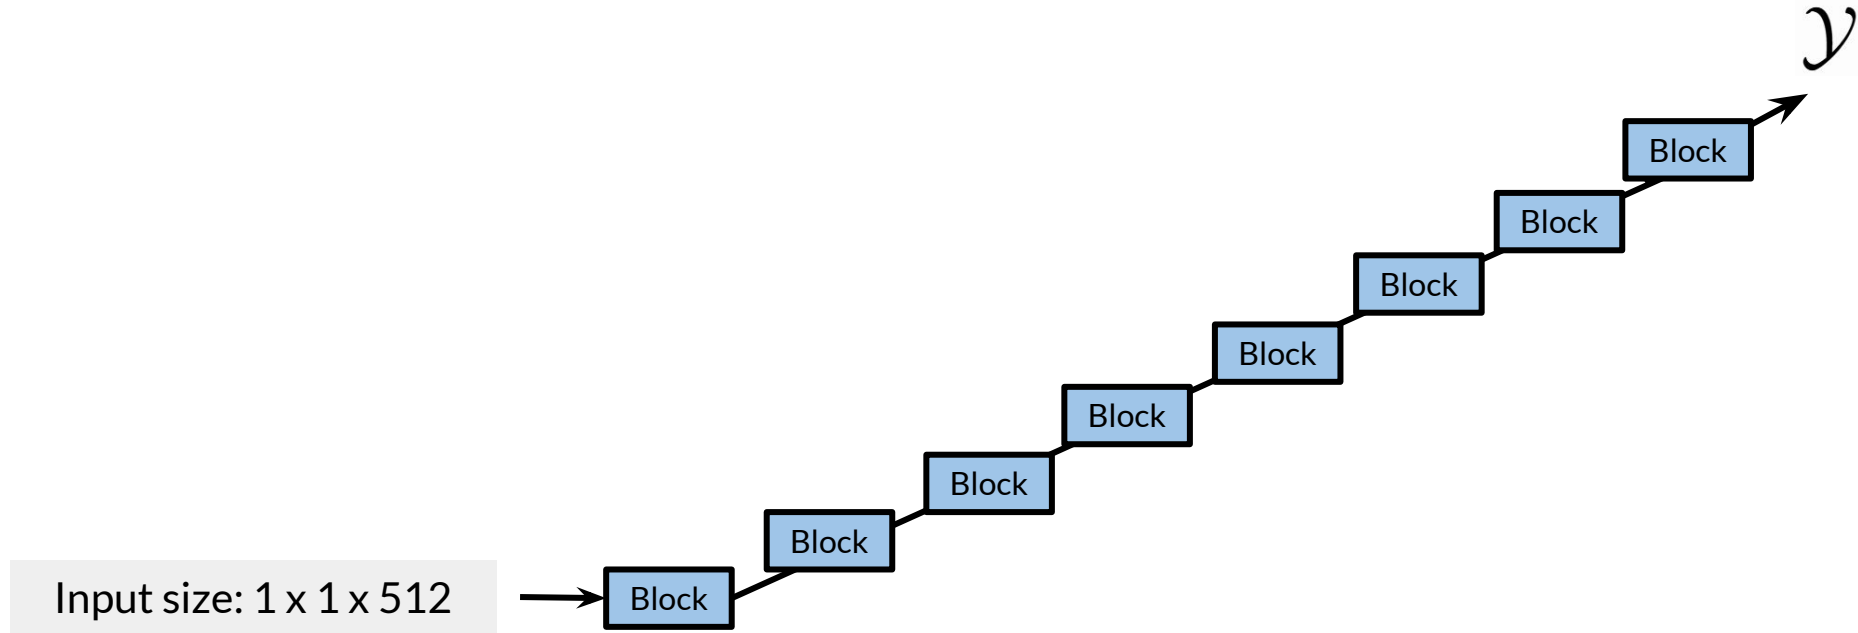
Pix2Pix Encoder



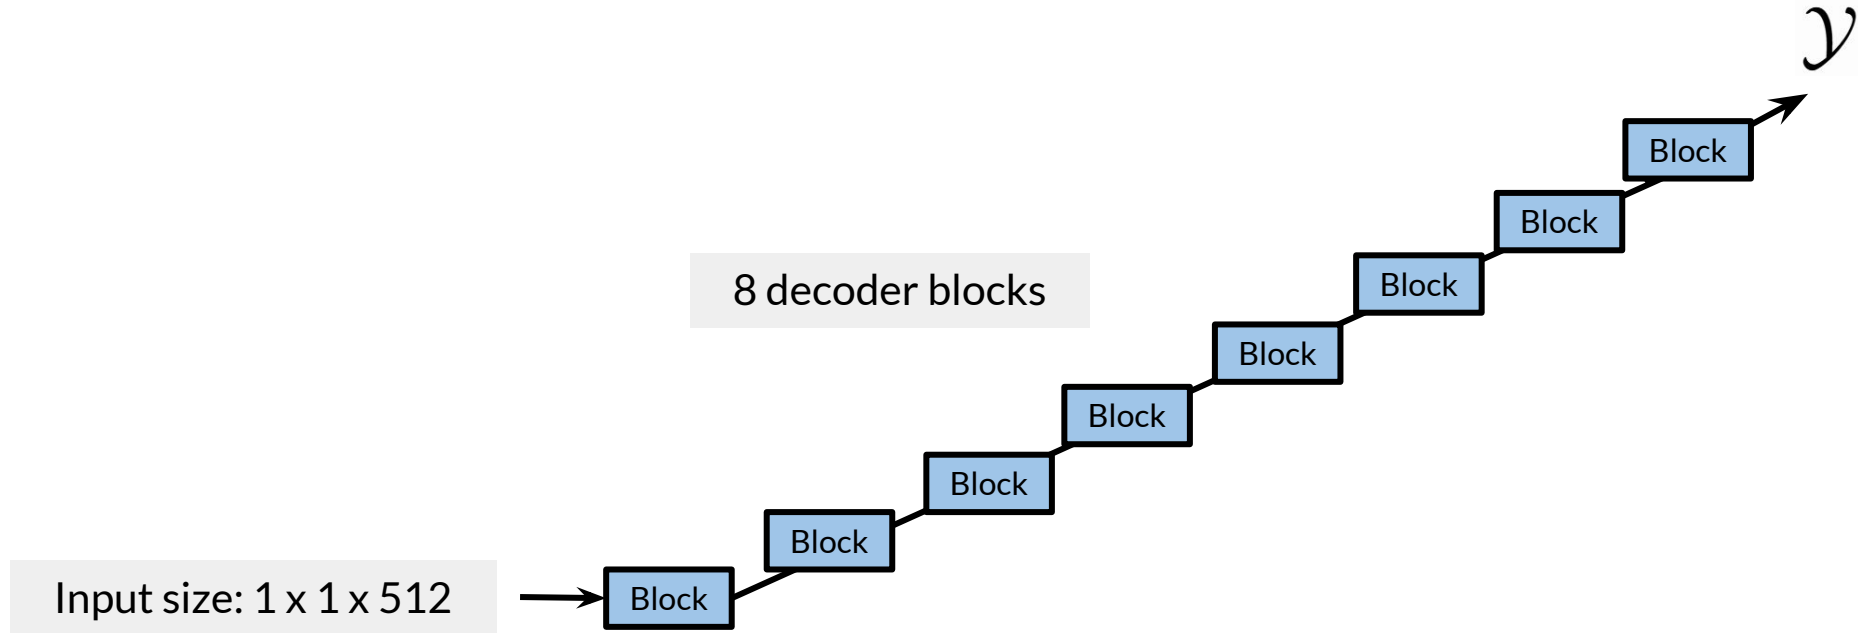
Pix2Pix Encoder



Pix2Pix Decoder



Pix2Pix Decoder



Pix2Pix Decoder

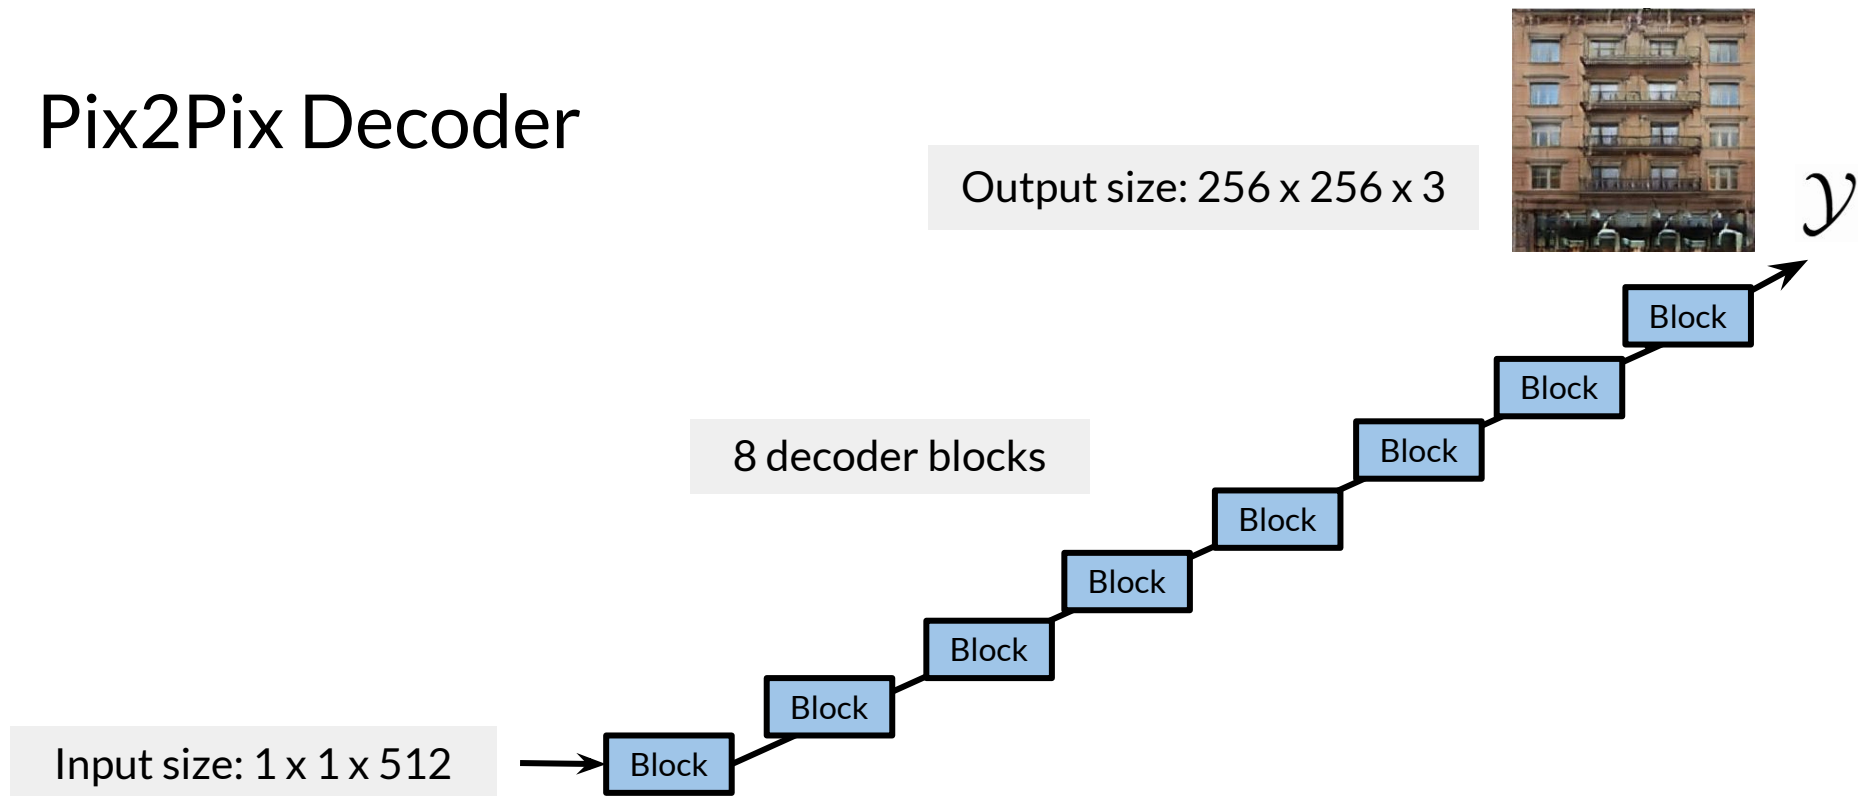


Image available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix Decoder

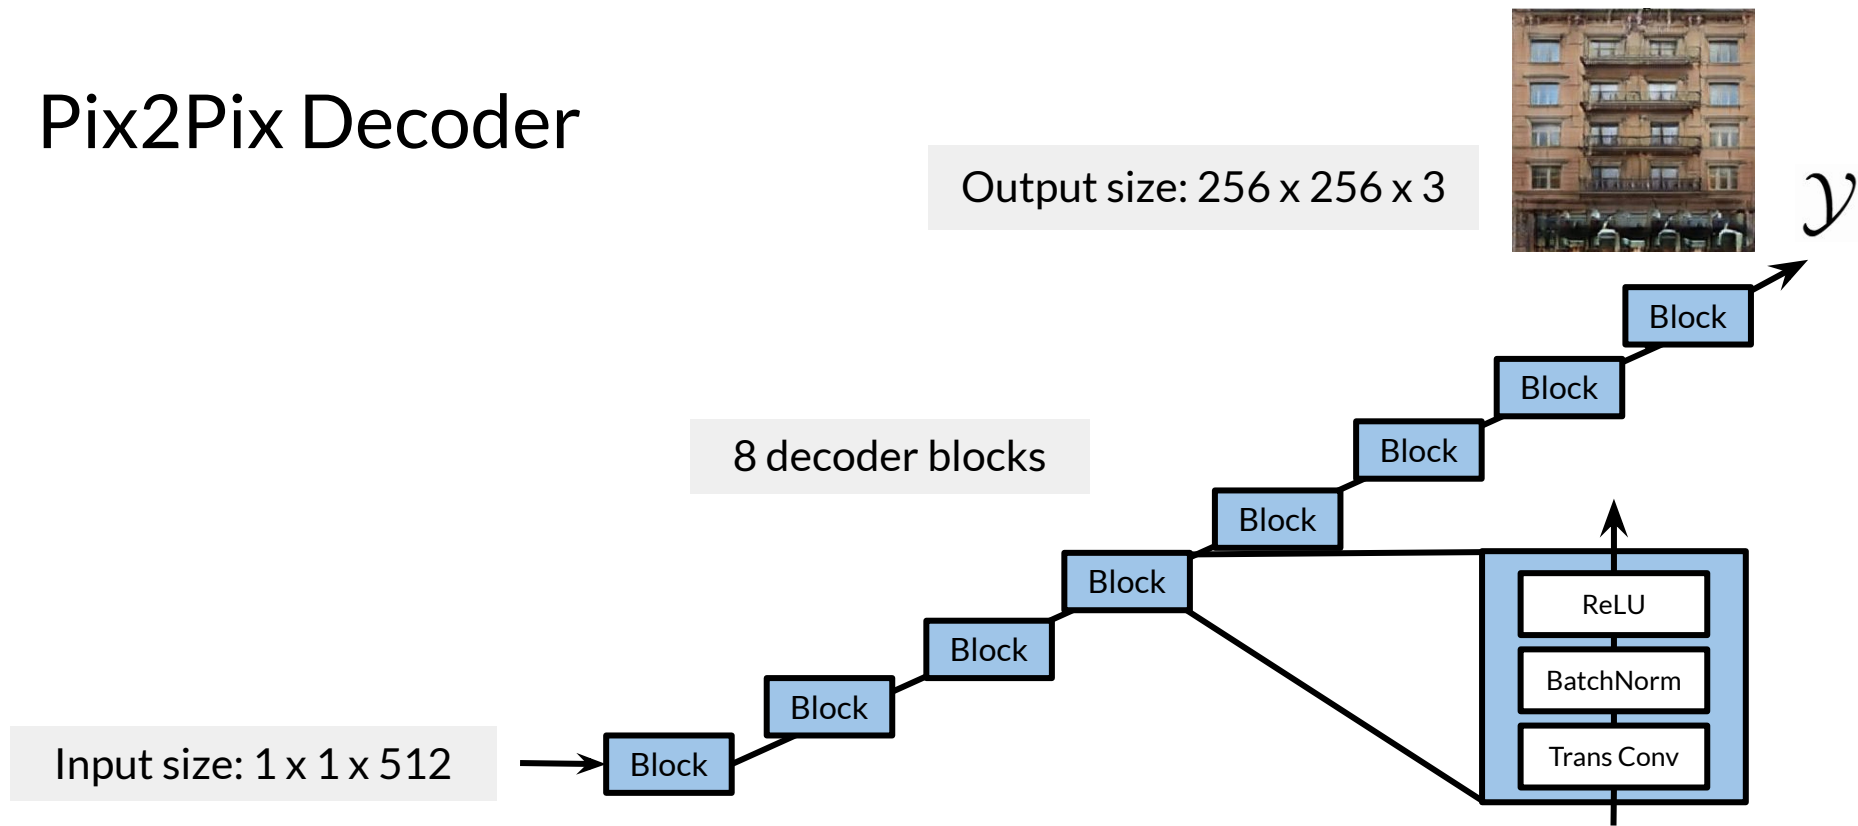



Image available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix Decoder

Output size: 256 x 256 x 3



y

8 decoder blocks

Dropout in some decoder blocks
adds noise to the network

Input size: 1 x 1 x 512


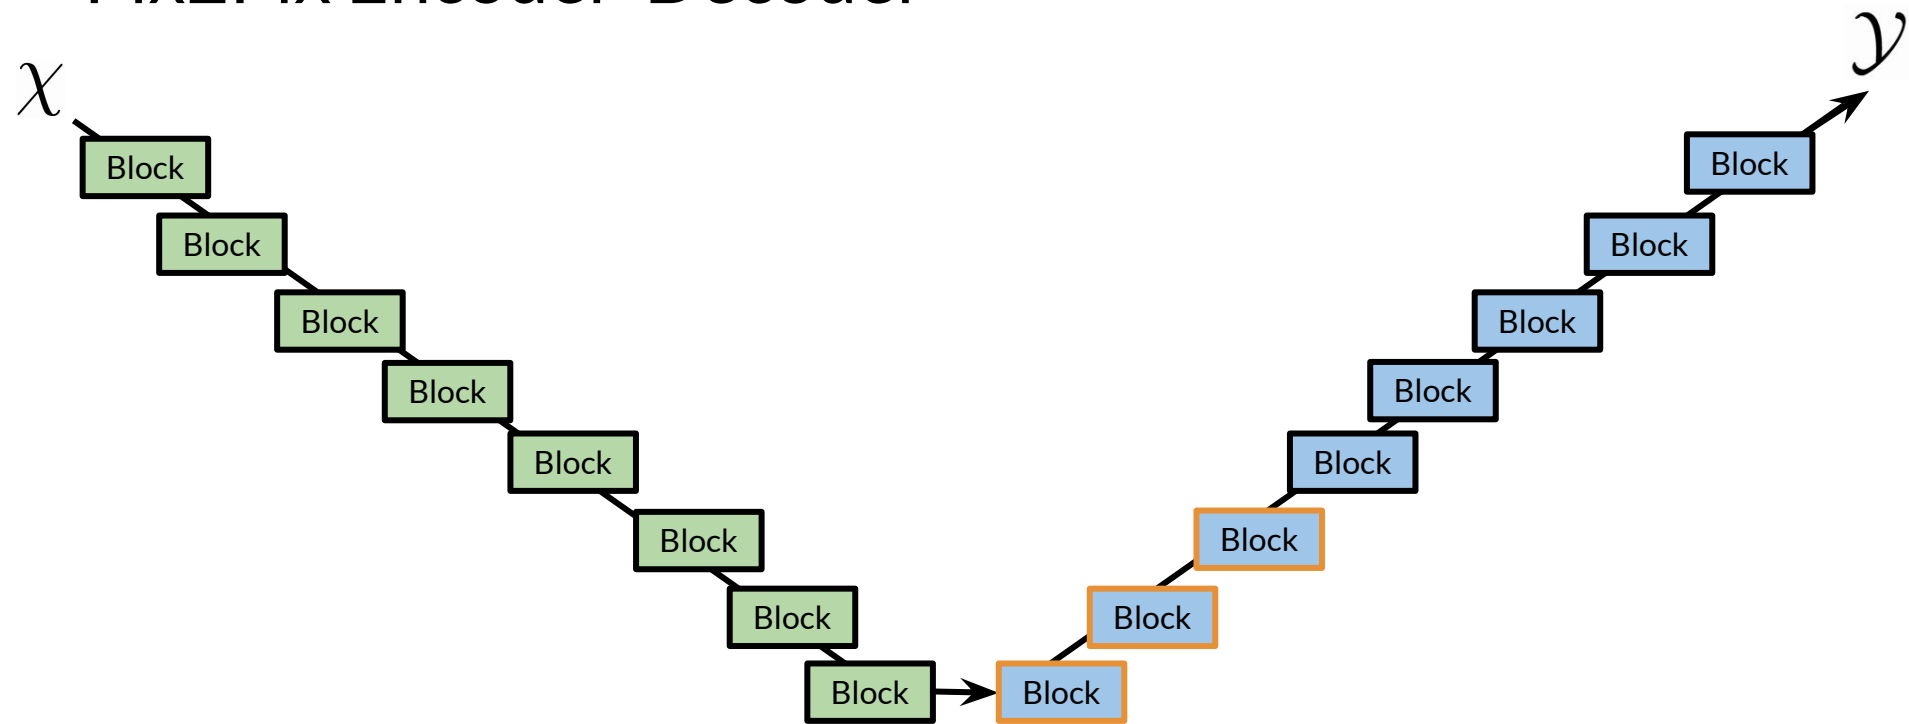
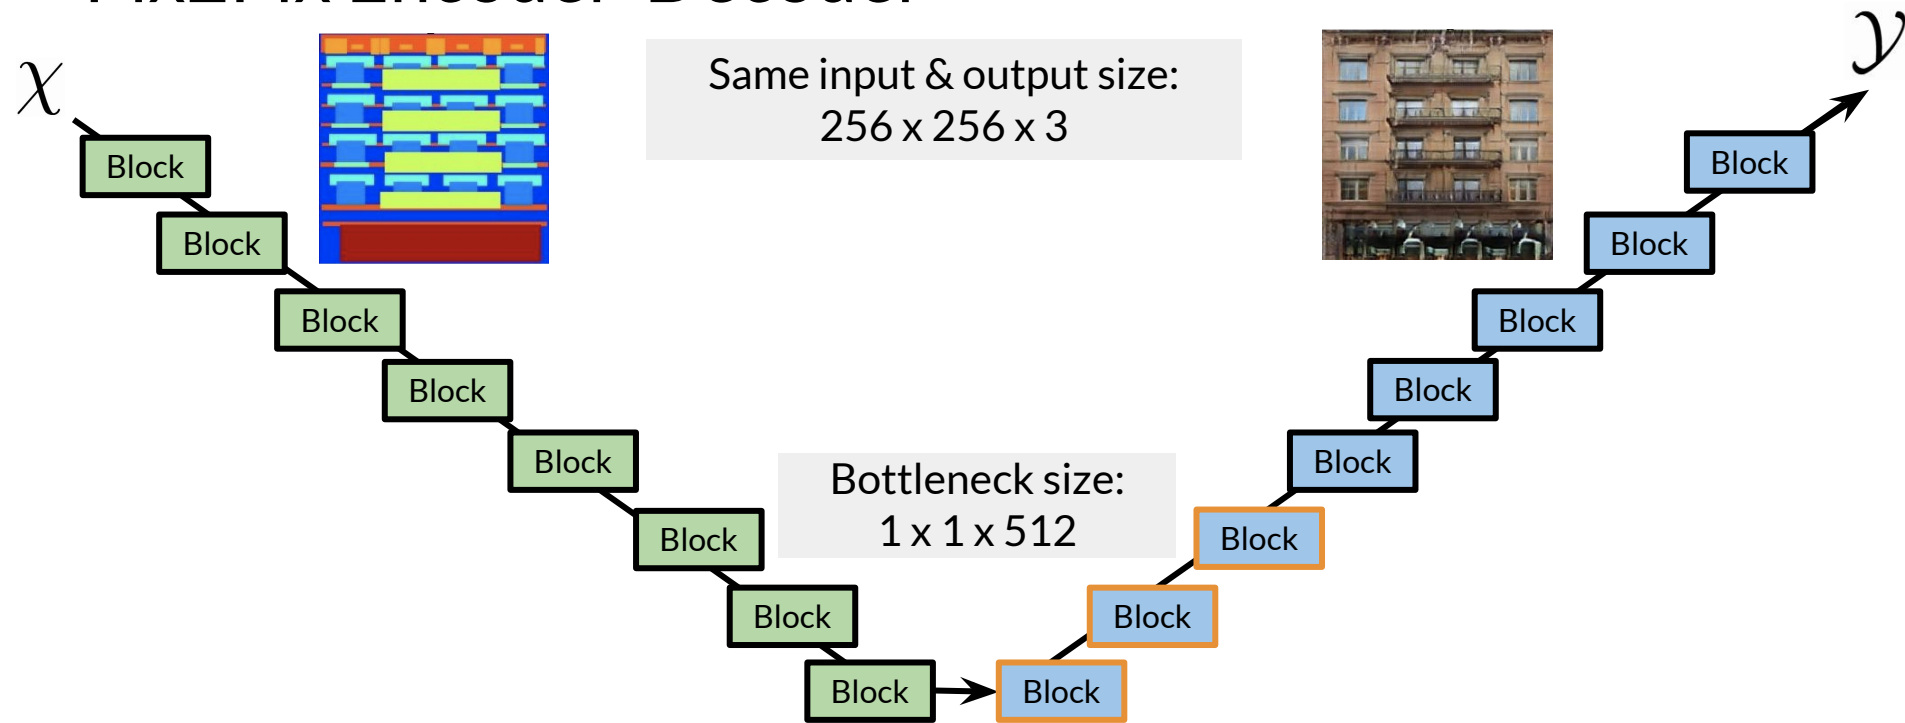


Image available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix Encoder-Decoder

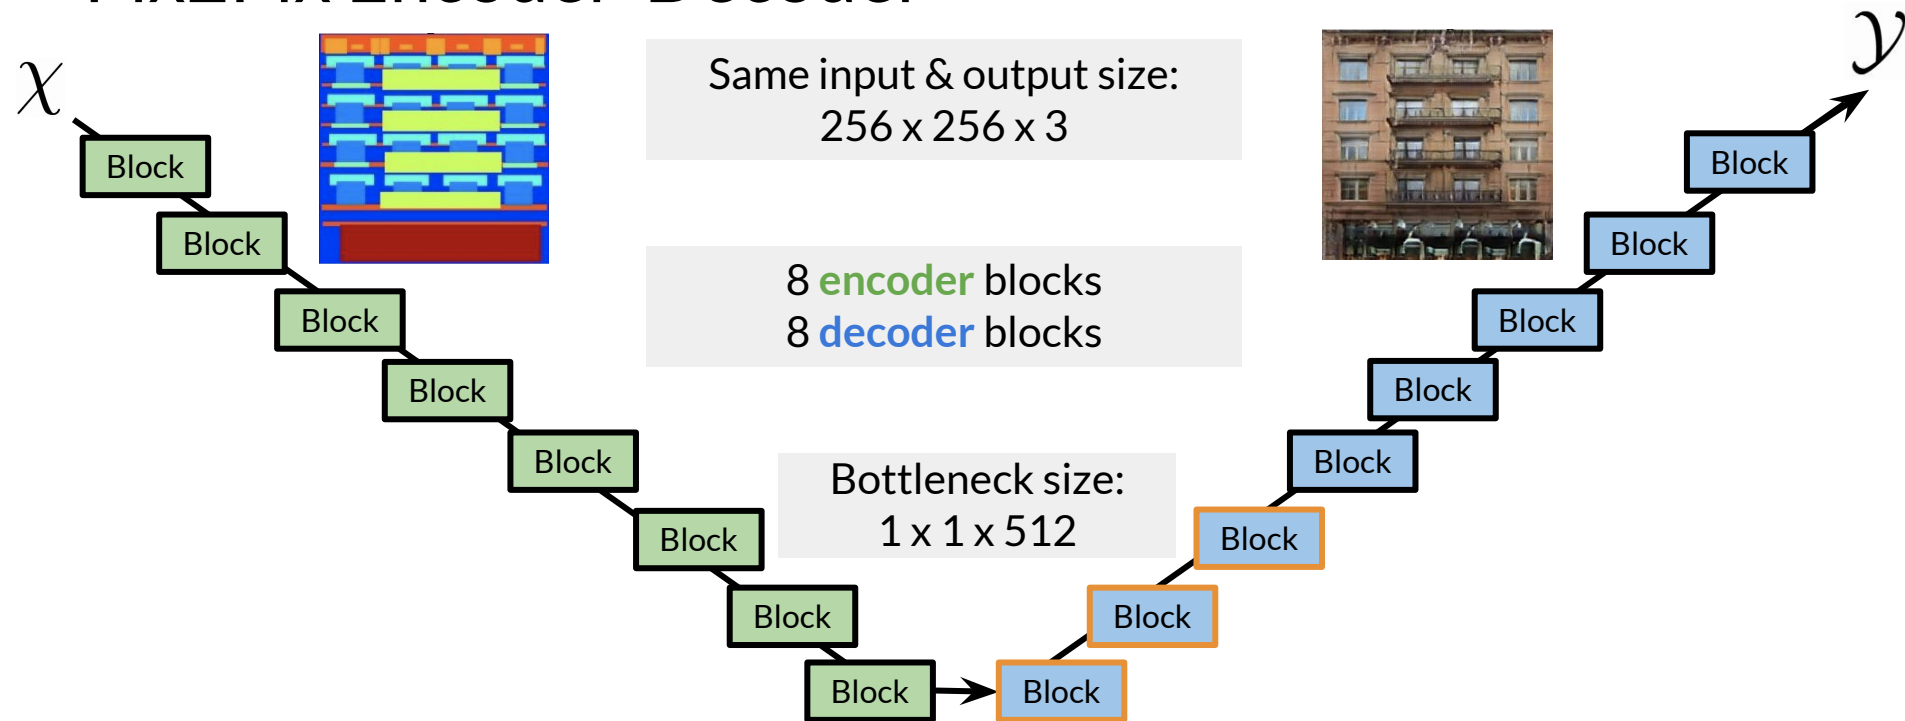


Pix2Pix Encoder-Decoder



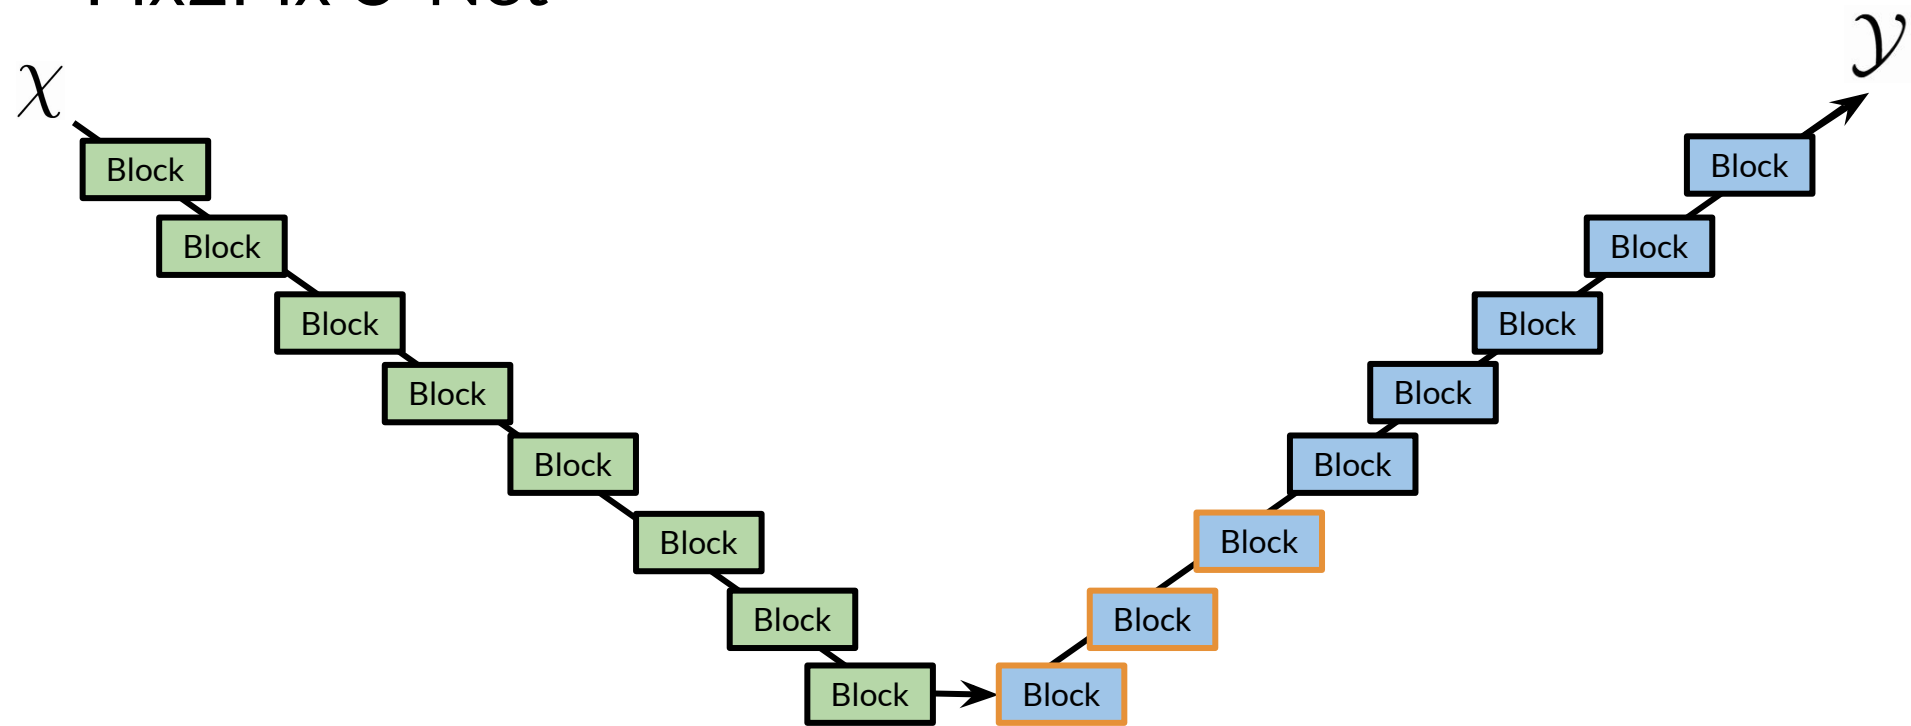
Images available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix Encoder-Decoder

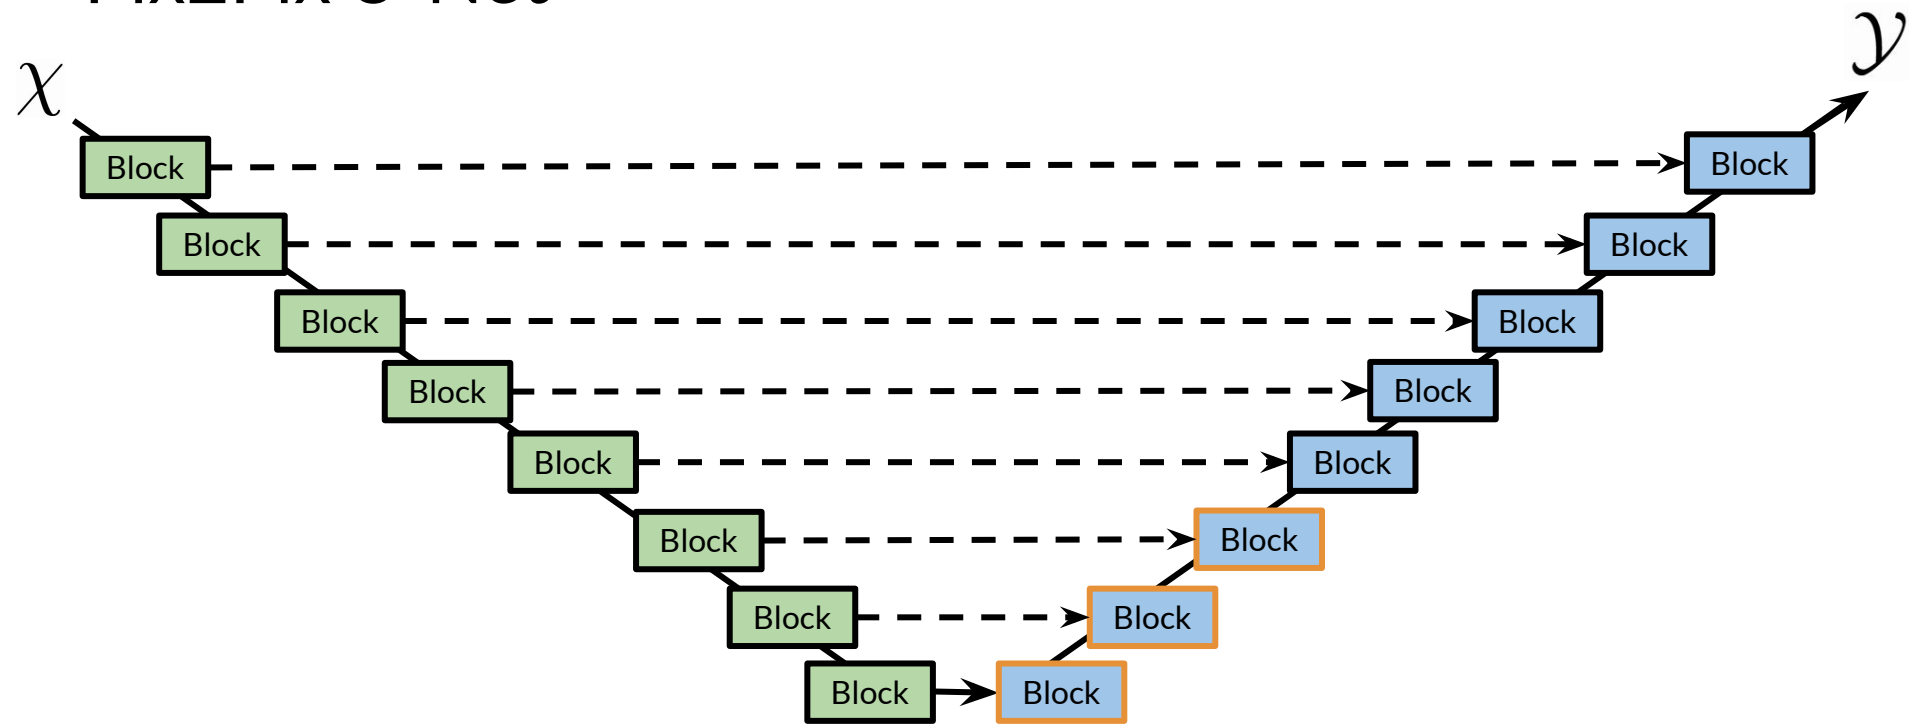


Images available from: <https://arxiv.org/abs/1611.07004>

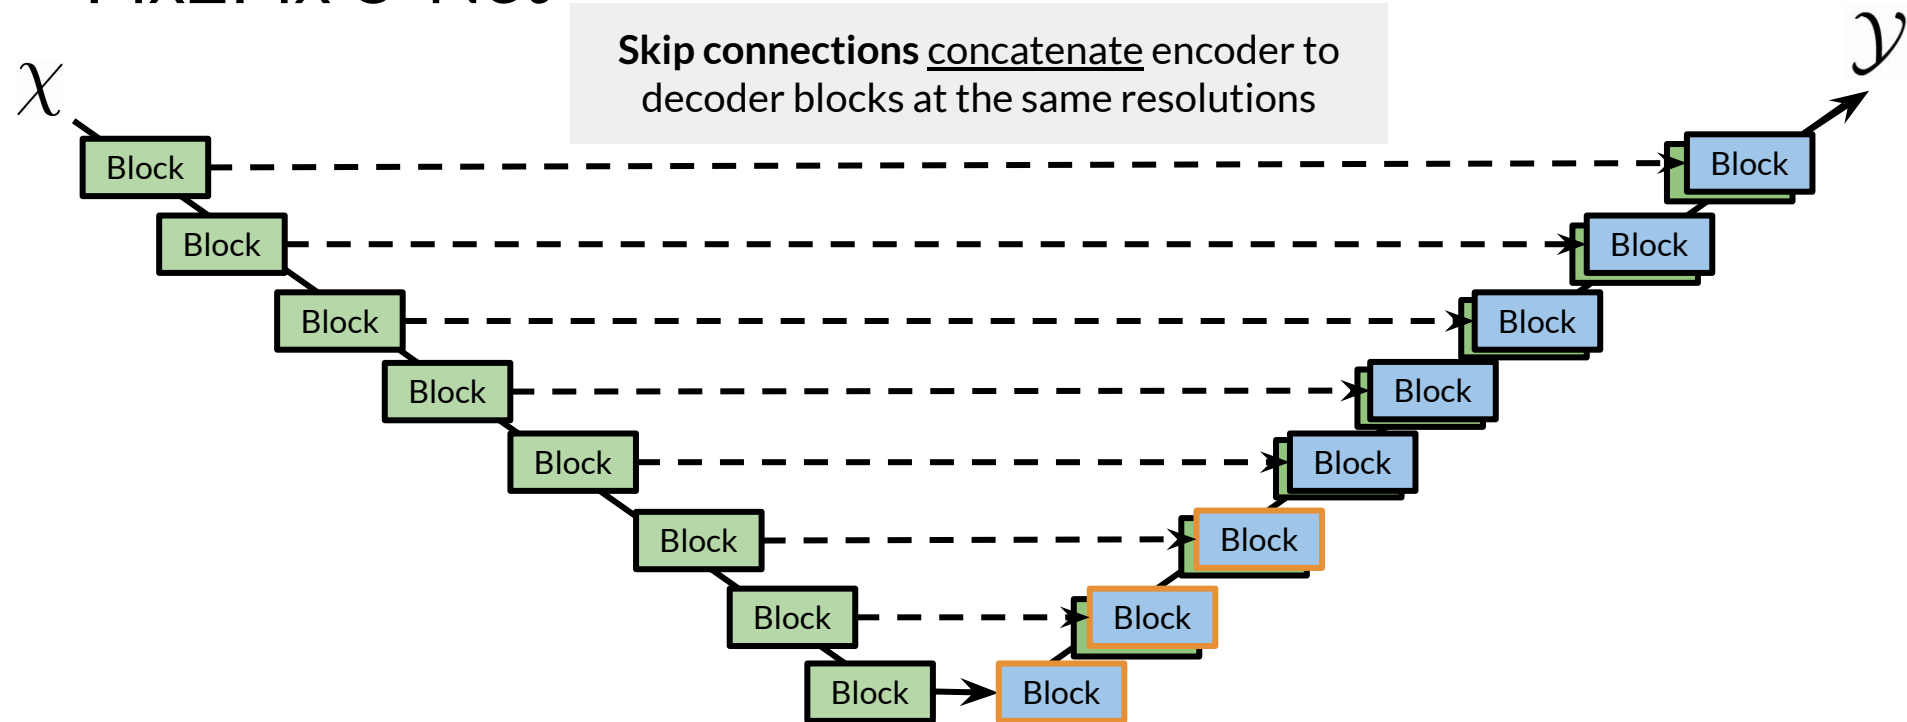
Pix2Pix U-Net



Pix2Pix U-Net




Pix2Pix U-Net



Summary

- Pix2Pix's generator is a U-Net
- U-Net is an encoder-decoder, with same-size inputs and outputs
- U-Net uses skip connections
 - Skip connections help the decoder learn details from the encoder directly
 - Skip connections the encoder learn from more gradients flowing from decoder






deeplearning.ai

Pix2Pix: Pixel Distance Loss Term

Outline

- Regularization and additional loss term
- Encourage pixel distance between generated and real outputs
- Additional loss term for Pix2Pix generator



Additional Loss Term

$$\min_g \max_c \text{Adversarial Loss} + \lambda * \text{Other loss term}$$

Additional Loss Term

$$\min_g \max_c \text{Adversarial Loss} + \lambda * \text{Pixel loss term}$$

Pixel Distance Loss Term

$$\sum_{i=1}^n$$



Generated output

—




Real output

Available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix Generator Loss

BCE Loss +

$$\lambda \sum_{i=1}^n$$



-



Available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix Generator Loss

$$\text{BCE Loss} + \lambda \sum_{i=1}^n \left| \text{generated_output} - \text{real_output} \right|$$

Summary

- Pix2Pix adds a Pixel Distance Loss term to the generator loss function
- This loss term calculates the difference between the fake and the real target outputs
- Softly encourages the generator with this additional supervision
 - The target output labels are the supervision
 - Generator essentially “sees” these labels






deeplearning.ai


Pix2Pix: Putting It All Together

Outline

- Put the Pix2Pix architecture together!
 - U-Net generator
 - Pixel Distance Loss term
 - PatchGAN discriminator




Pix2Pix



Real input


Image available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix



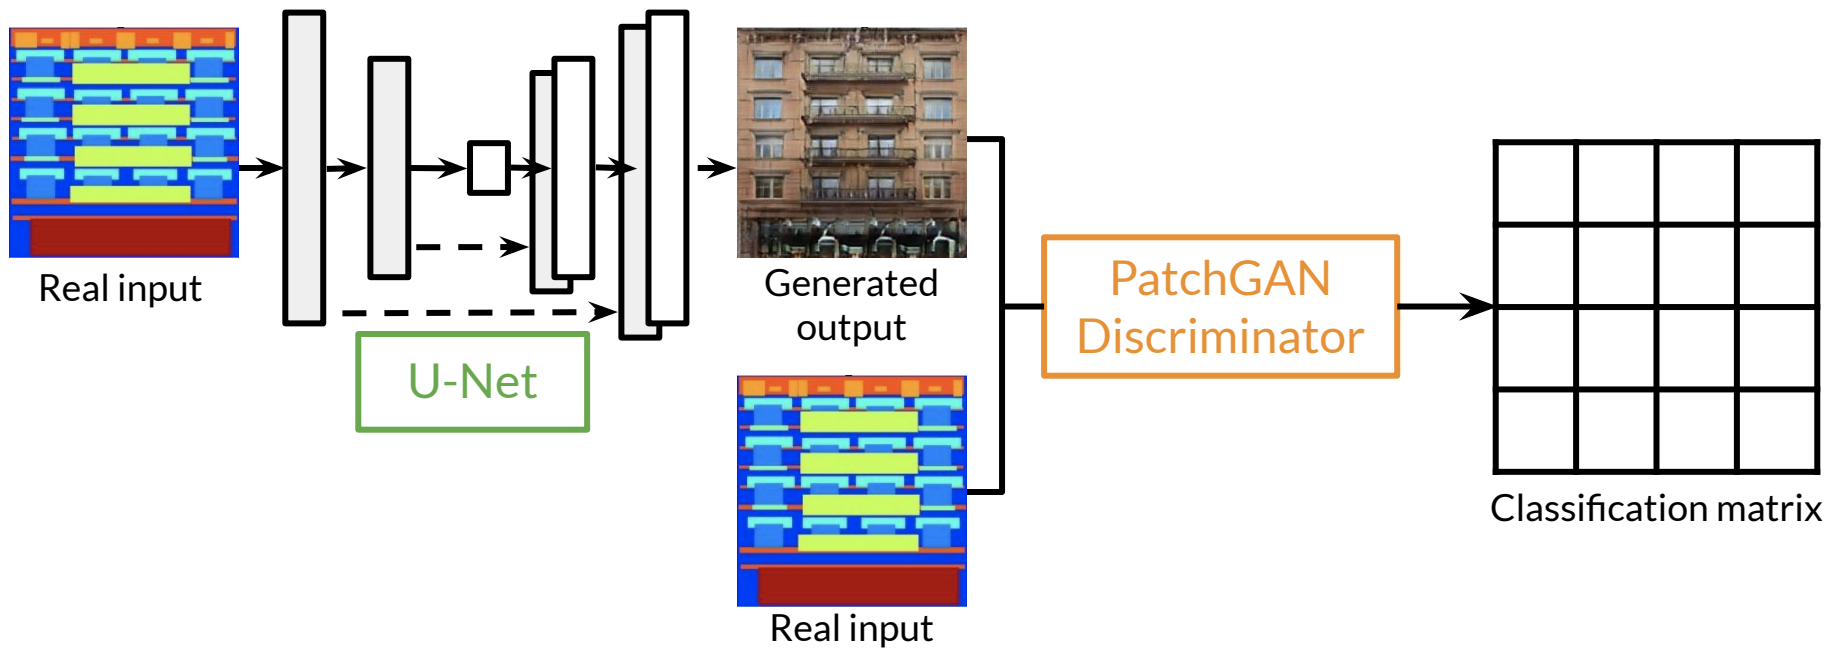
Images available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix




Images available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix




Images available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix: Discriminator Loss




Images available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix: Discriminator Loss




Images available from: <https://arxiv.org/abs/1611.07004>

Pix2Pix: Generator Loss



Images available from: <https://arxiv.org/abs/1611.07004>


Pix2Pix: Generator Loss



Images available from: <https://arxiv.org/abs/1611.07004>

Summary

- U-Net generator: image \rightarrow image
- PatchGAN discriminator
 - Inputs input image and paired output (either real target or fake)
 - Outputs classification matrix
- Generator loss has a regularization term






deeplearning.ai

Pix2Pix Advancements

Outline

- Improvements and extensions of Pix2Pix for paired image-to-image translation
 - Higher resolution images
 - Image editing



Pix2PixHD



Available from: <https://github.com/NVIDIA/pix2pixHD>

GauGAN



Available from: <https://blogs.nvidia.com/blog/2019/03/18/gaugan-photorealistic-landscapes-nvidia-research/>

Summary

- Pix2PixHD and GauGAN are successors of Pix2Pix
- They are designed for higher resolution images
- They highlight opportunities for image editing using paired image-to-image translation
 - Pix2Pix can do this too, of course!

